

# Spring Batch Technology Preview + Q&A

Kansas City Spring User Group  
August 5, 2015

David Pitt, Managing Partner  
[dpitt@keyholesoftware.com](mailto:dpitt@keyholesoftware.com)



Presentation & Demo © Keyhole Software 2015

# Batch Processing

- Emphasis seems to be on SOA and real time integration
- Still lots of mission critical apps processing billions of transactions per day in “batch mode”
- Many interfaces are still “file-based” batch processes
- Lack of architecture and framework specs have lead to expensive one-off and in-house custom batch implementations

How many times have you created a Java main(..) method, kicked off from a cron job to process an import file....?

# Spring Batch – What Is It?

- ETL Batch processing framework for Java
- Jobs built using POJOs and Spring Dependency Injection(DI) framework
- Provides Execution Runtimes for Jobs
- Ability to Process Large Amounts of Data
- Transaction Integrity and Restartable

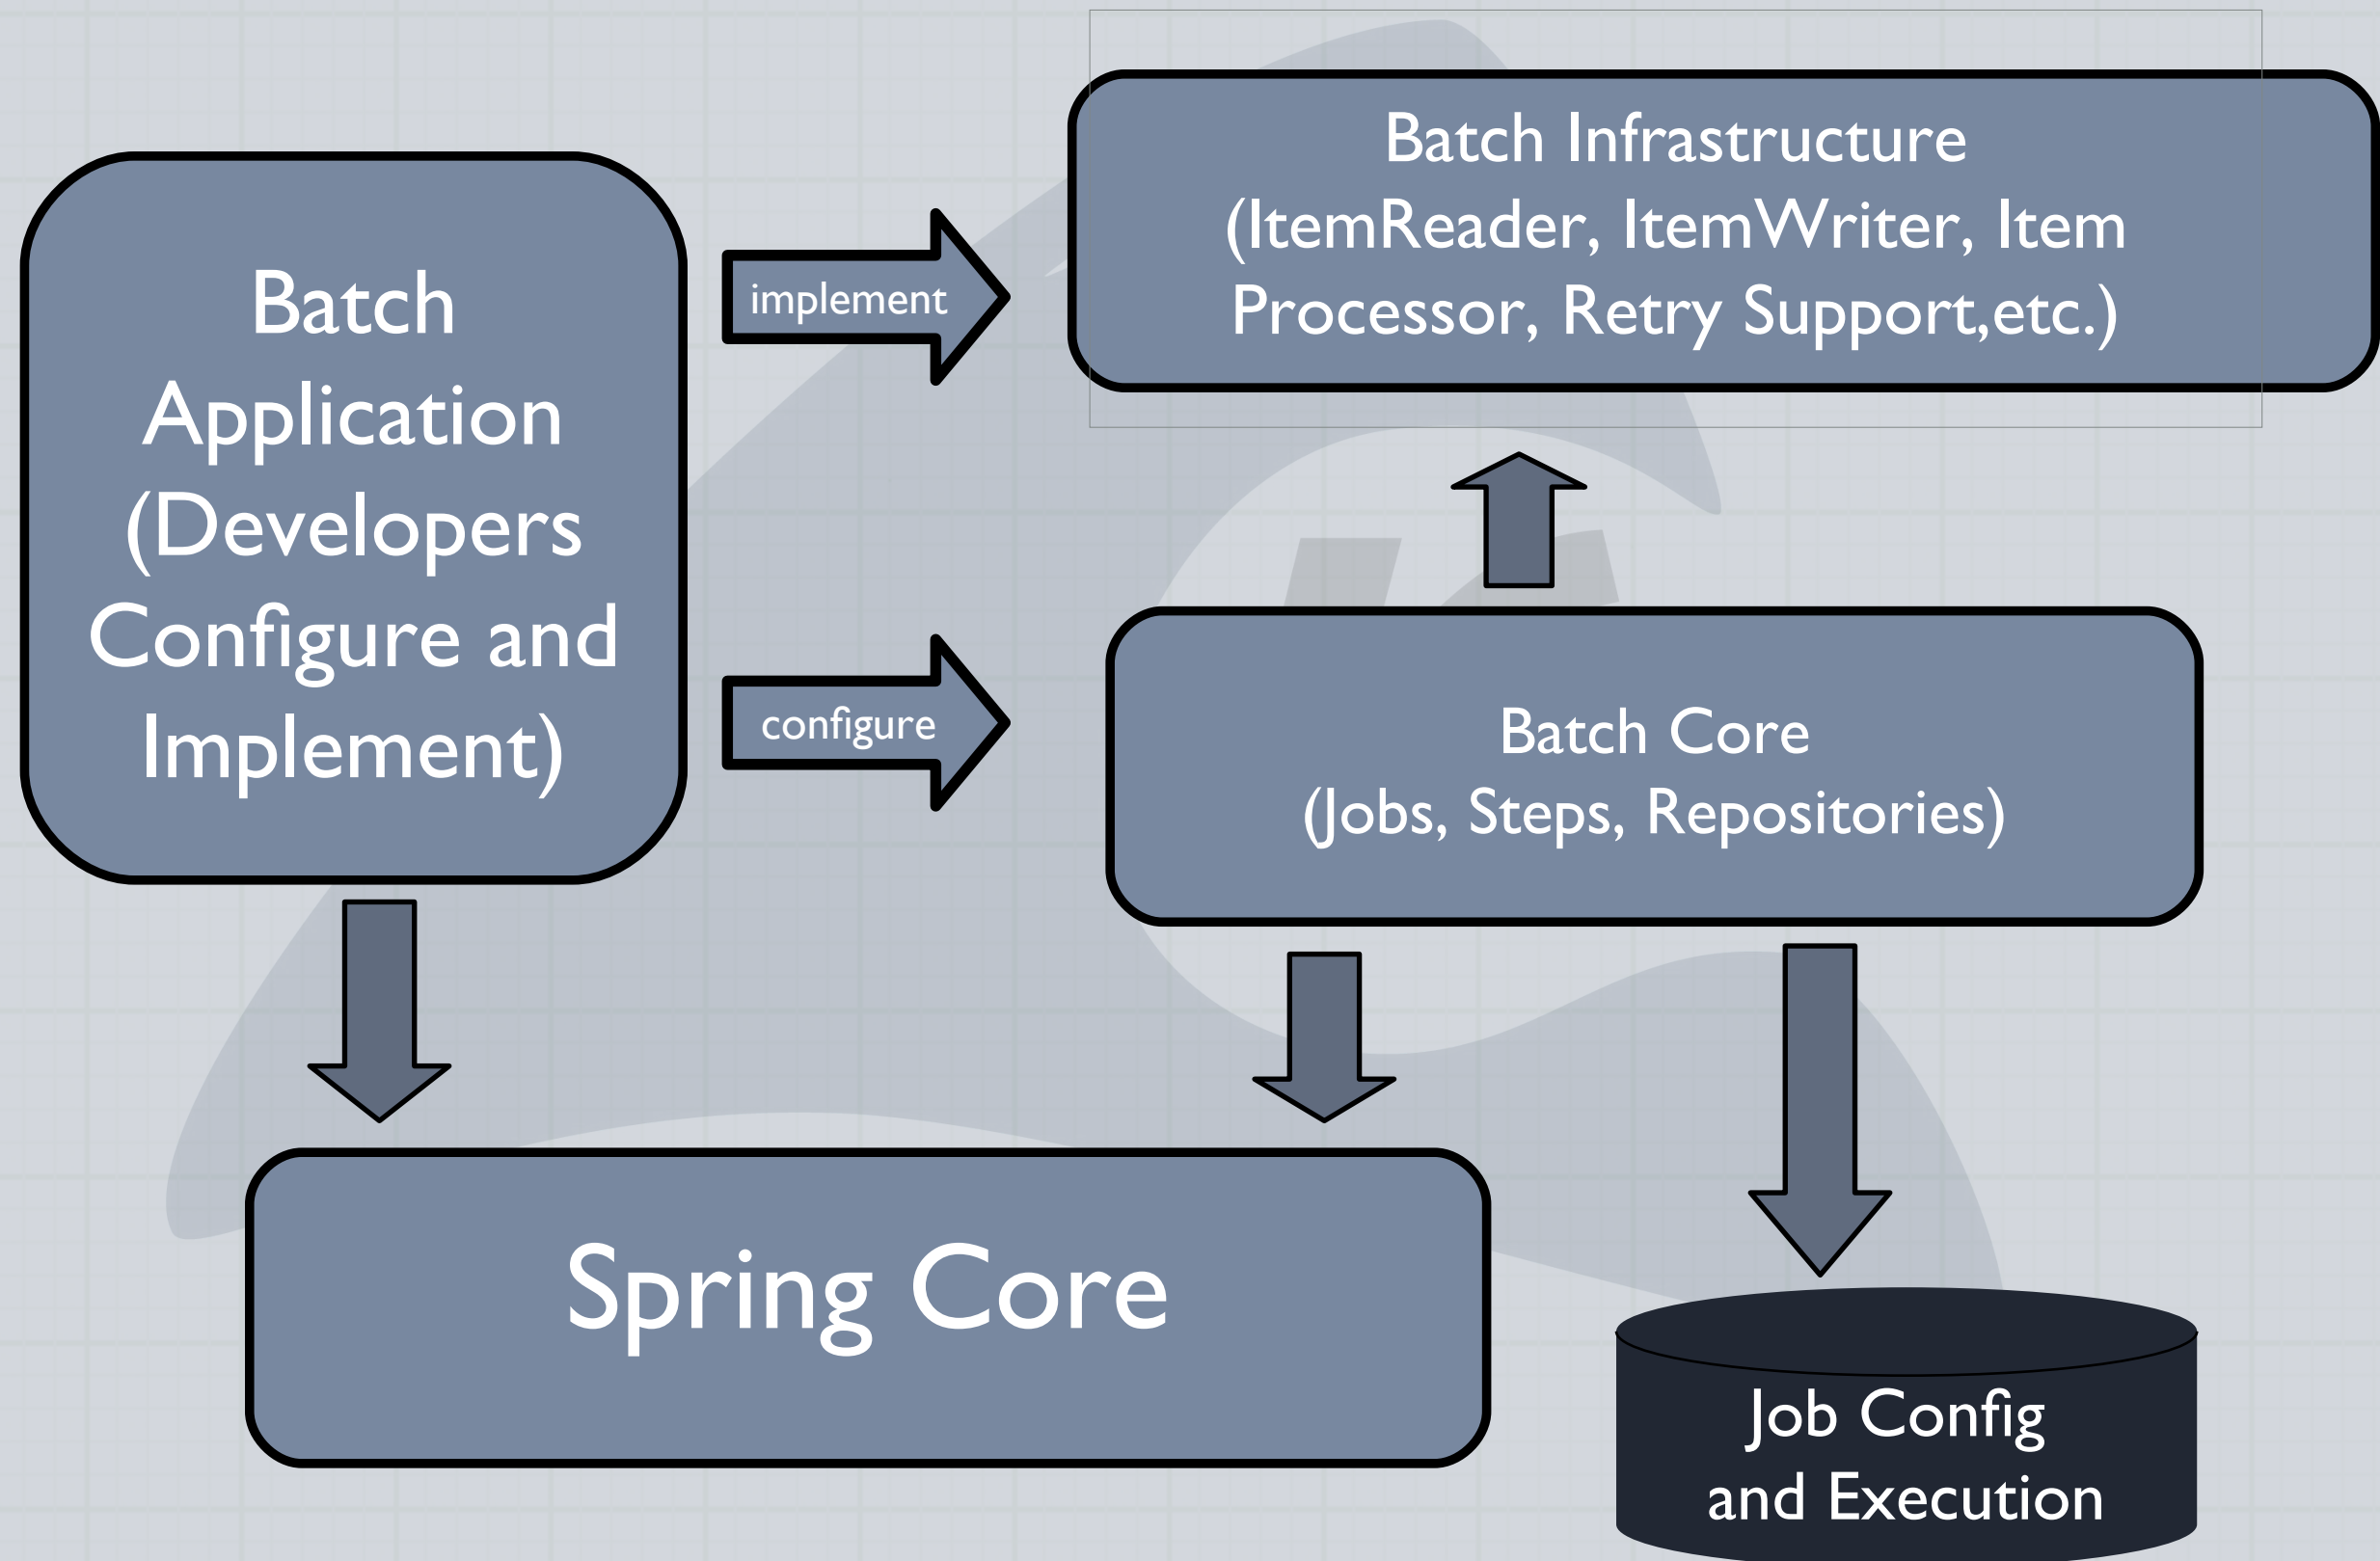
# Where It Came From

- Collaboration between Accenture and Spring (Vmware), now Pivotal
- Accenture had many years experience with numerous enterprise batch processing environments and solutions
- Spring simplifies and modularizes JEE application architecture frameworks via dependency injection



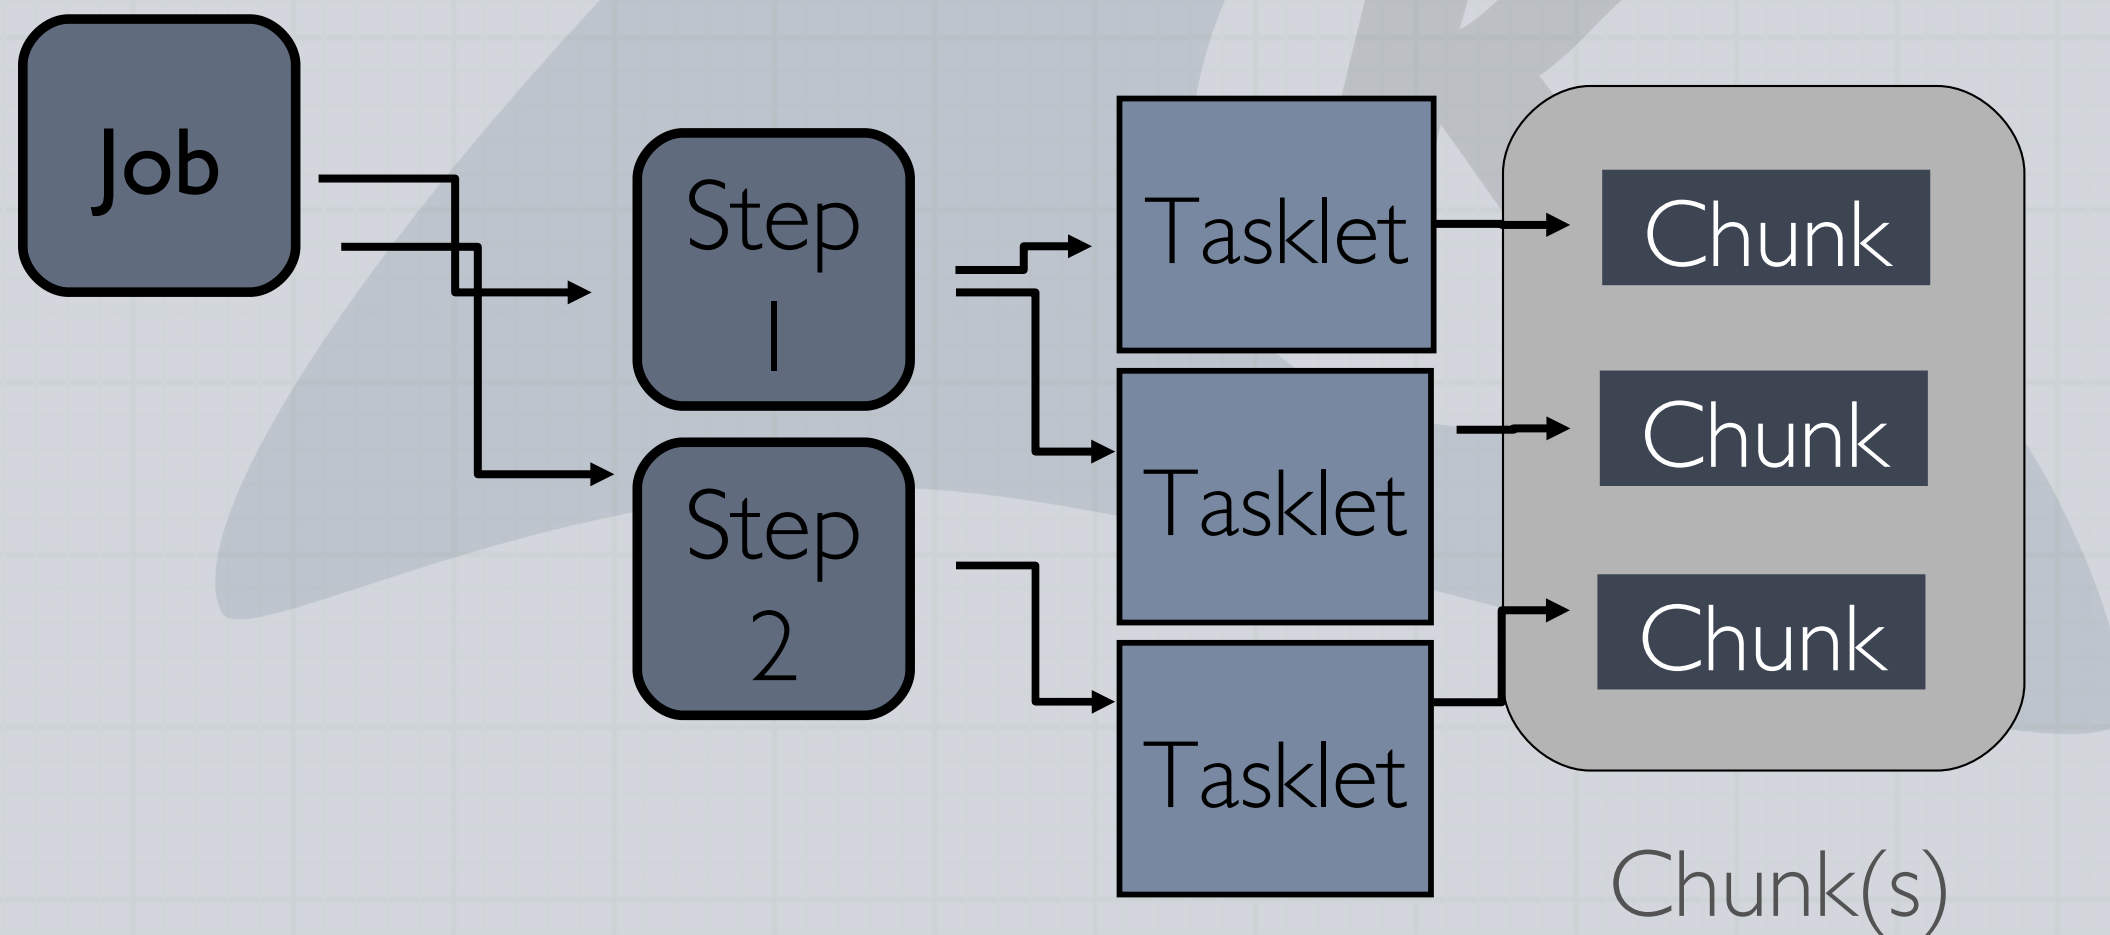
Pivotal™

# Architecture



# Concepts and Features

- Borrows from Traditional Batch Processing Concepts
  - Jobs are made up from steps that perform tasks
  - Run in a single VM?





# Job Definition

- Spring DI used to configure (WIRE UP) classes configuration
- Elements to “Wire UP” in application Context

- JobOperator
- JobRepository
- JobLauncher
- Job
- Steps
- Tasklet
- Chunks (Reader/Processor/Writer)
- Standard Stuff (Properties, DBconnections, tx manager, etc)

Spring Core  
Provides Impl

Developer  
Configures

Implement

# Job Repository

- Governs CRUD for domain objects during job and step execution
- Repositories can be a database or in memory (default)

```
<job-repository id="jobRepository"  
    data-source="dataSource"  
    transaction-manager="transactionManager"  
    isolation-level-for-create="SERIALIZABLE"  
    table-prefix="BATCH_" max-varchar-length="1000" />
```

Job  
Launcher

Job  
Instance

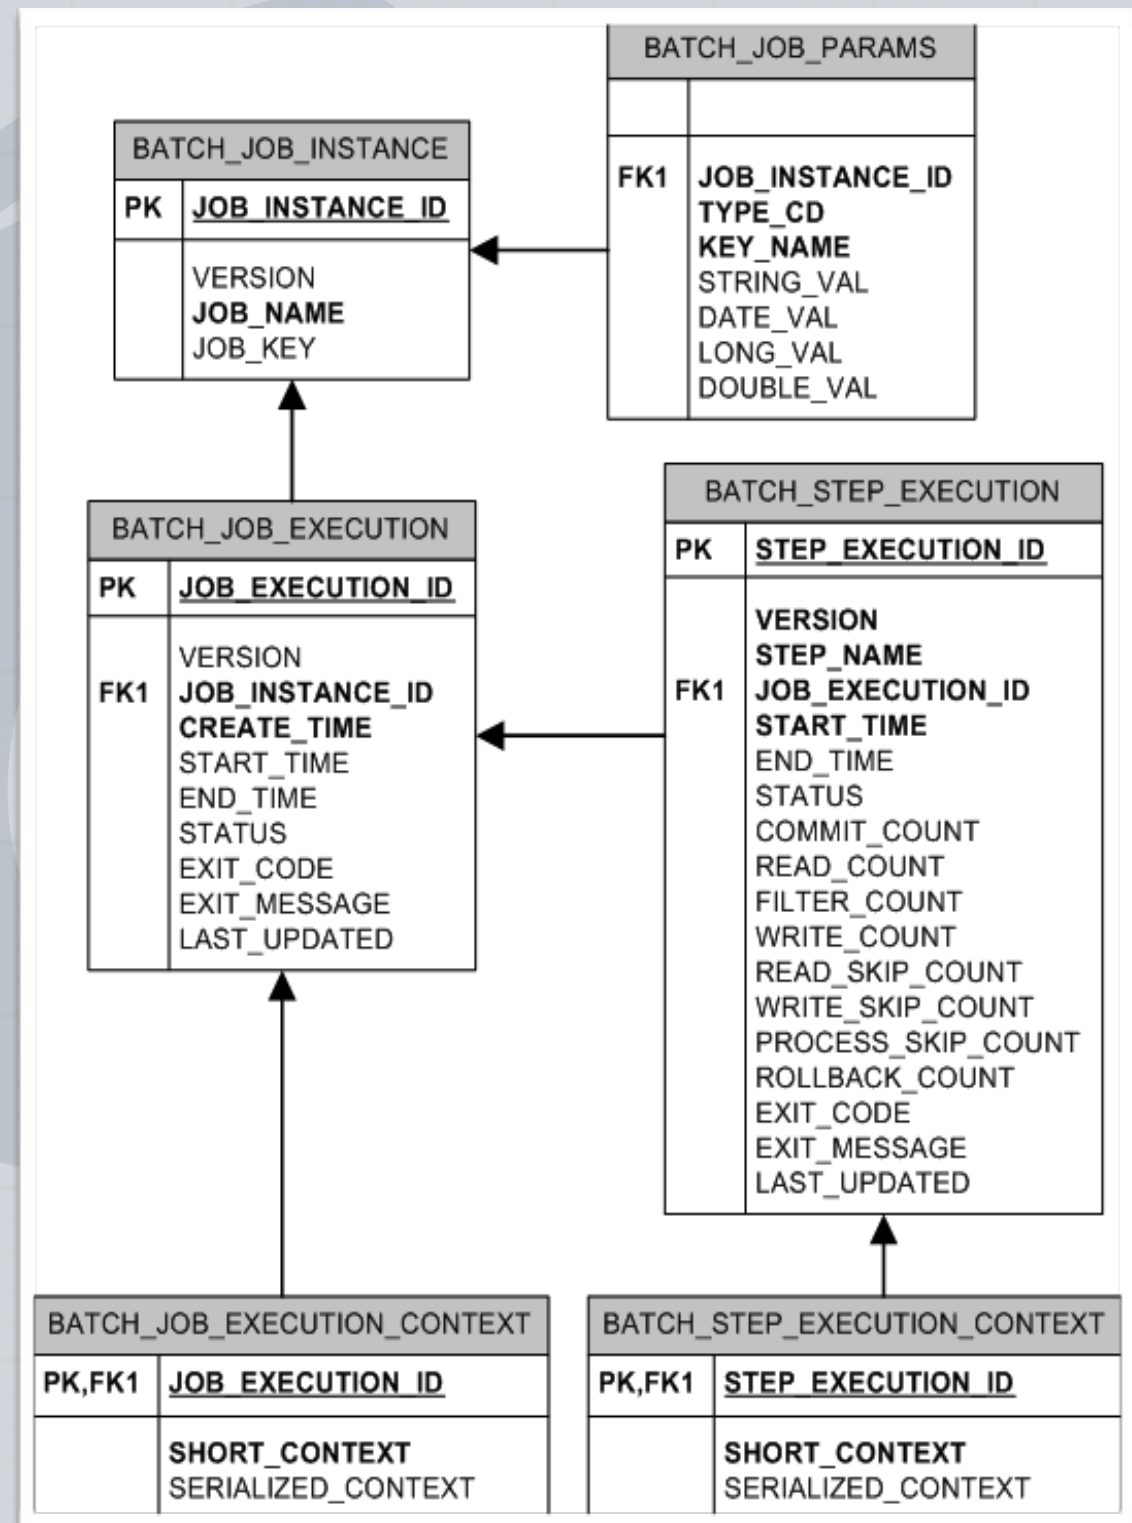
Step

JobRepository (DB2, Oracle, any Relational DB)



# Batch META-DATA

## Relational Meta Data Supporting Spring Batch



# Job

- Encapsulates entire Job Process steps
- Configured in Spring job context(s) files
- Or, Java Config

```
<job id="TimesheetInvoicing" restartable="false">
  <step id="importTimesheets" next="invoiceDetail">
    <tasklet transaction-manager="txManager" >
      <chunk reader="timesheetFileItemReader"
        writer="timesheetWriter"
        commit-interval="3">
        </chunk>
      </tasklet>
    </step>
    <step id="invoiceDetail" next="outputPdfReport">
      <tasklet transaction-manager="txManager" >
        <chunk reader="timesheetReader"
          processor="createInvoice"
          writer="writeInvoice"
          commit-interval="3">
          </chunk>
        </tasklet>
      </step>
      <step id="outputPdfReport">
        <tasklet transaction-manager="txManager" >
          <chunk reader="reportReader"
            writer="reportWriter"
            commit-interval="3">
            </chunk>
          </tasklet>
        </step>
      </job>
```

# Job Java Config

```
@Configuration
@EnableBatchProcessing
public class TickerPriceConversionConfig {

    @Autowired
    private JobBuilderFactory jobs;

    @Autowired
    private StepBuilderFactory steps;

    @Bean
    public ItemReader<TickerData> reader() throws MalformedURLException {
        FlatFileItemReader<TickerData> reader = new FlatFileItemReader<TickerData>();
        reader.setResource(new UrlResource("http://finance.yahoo.com/d/quotes.csv?s=XOM+IBM+JNJ+MSFT&f=snd1o1p2"));
        reader.setLineMapper(new DefaultLineMapper<TickerData>() {{
            setLineTokenizer(new DelimitedLineTokenizer());
            setFieldSetMapper(new TickerFieldSetMapper());
        }});
        return reader;
    }

    @Bean
    public ItemProcessor<TickerData, TickerData> processor() {
        return new TickerPriceProcessor();
    }

    @Bean
    public ItemWriter<TickerData> writer() {
        return new LogItemWriter();
    }

    @Bean
    public Job TickerPriceConversion() throws MalformedURLException {
        return jobs.get("TickerPriceConversion").start(convertPrice()).build();
    }

    @Bean
    public Step convertPrice() throws MalformedURLException {
        return steps.get("convertPrice")
            .<TickerData, TickerData> chunk(5)
            .reader(reader())
            .processor(processor())
            .writer(writer())
            .build();
    }
}
```

# Steps

- Can be sequential or conditional

Step A

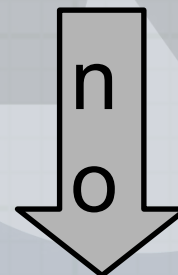


Step B

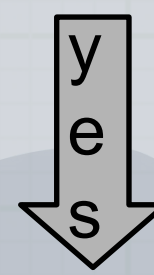


Step C

Step A



Step B



Step C

# Step Types

Spring Core Provides:

- Simple Step
- Fault Tolerant Step
- Skip and Retry Behavior
- Or, Implement Custom





# Step Definition

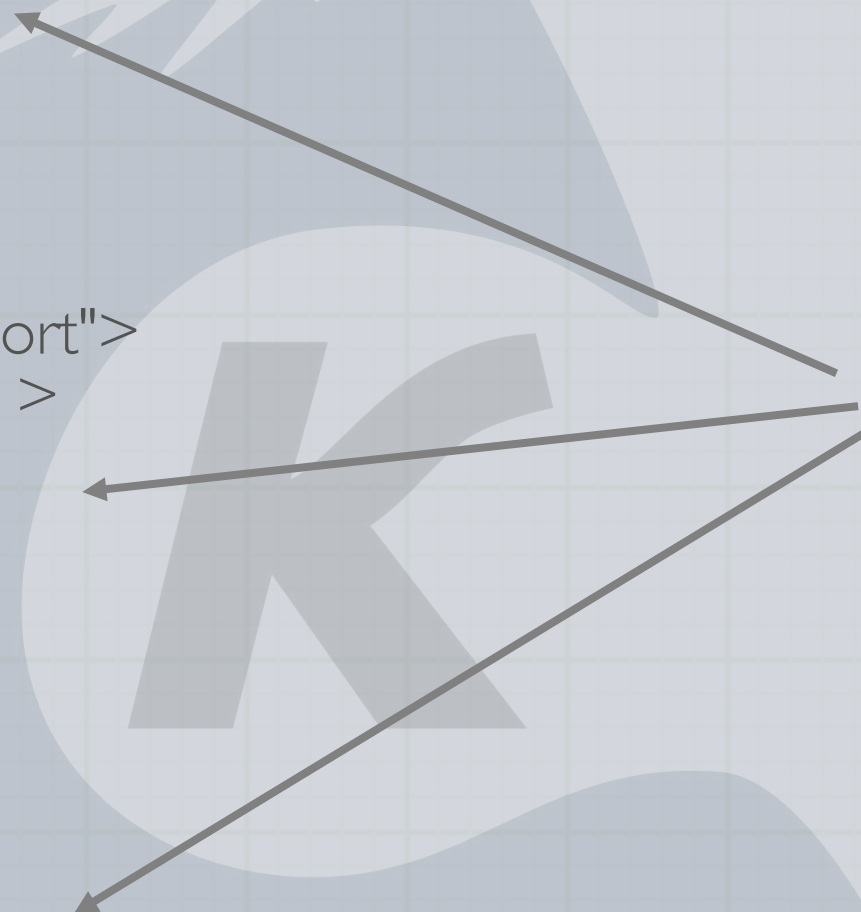
```
<job id="TimesheetInvoicing" restartable="false">  
<step id="importTimesheets" next="invoiceDetail">  
  <tasklet transaction-manager="txManager" >  
    <chunk reader="timesheetFileItemReader"
```

```
      writer="timesheetWriter"  
      commit-interval="3">  
    </chunk>  
  </tasklet>  
</step>
```

```
<step id="invoiceDetail" next="outputPdfReport">  
  <tasklet transaction-manager="txManager" >  
    <chunk reader="timesheetReader"  
      processor="createInvoice"  
        writer="writeInvoice"  
        commit-interval="3">  
      </chunk>  
    </tasklet>  
  </step>
```

```
<step id="outputPdfReport">  
  <tasklet transaction-manager="txManager" >  
    <chunk reader="reportReader"  
      writer="reportWriter"  
      commit-interval="3">  
    </chunk>  
  </tasklet>  
</step>  
</job>
```

Steps





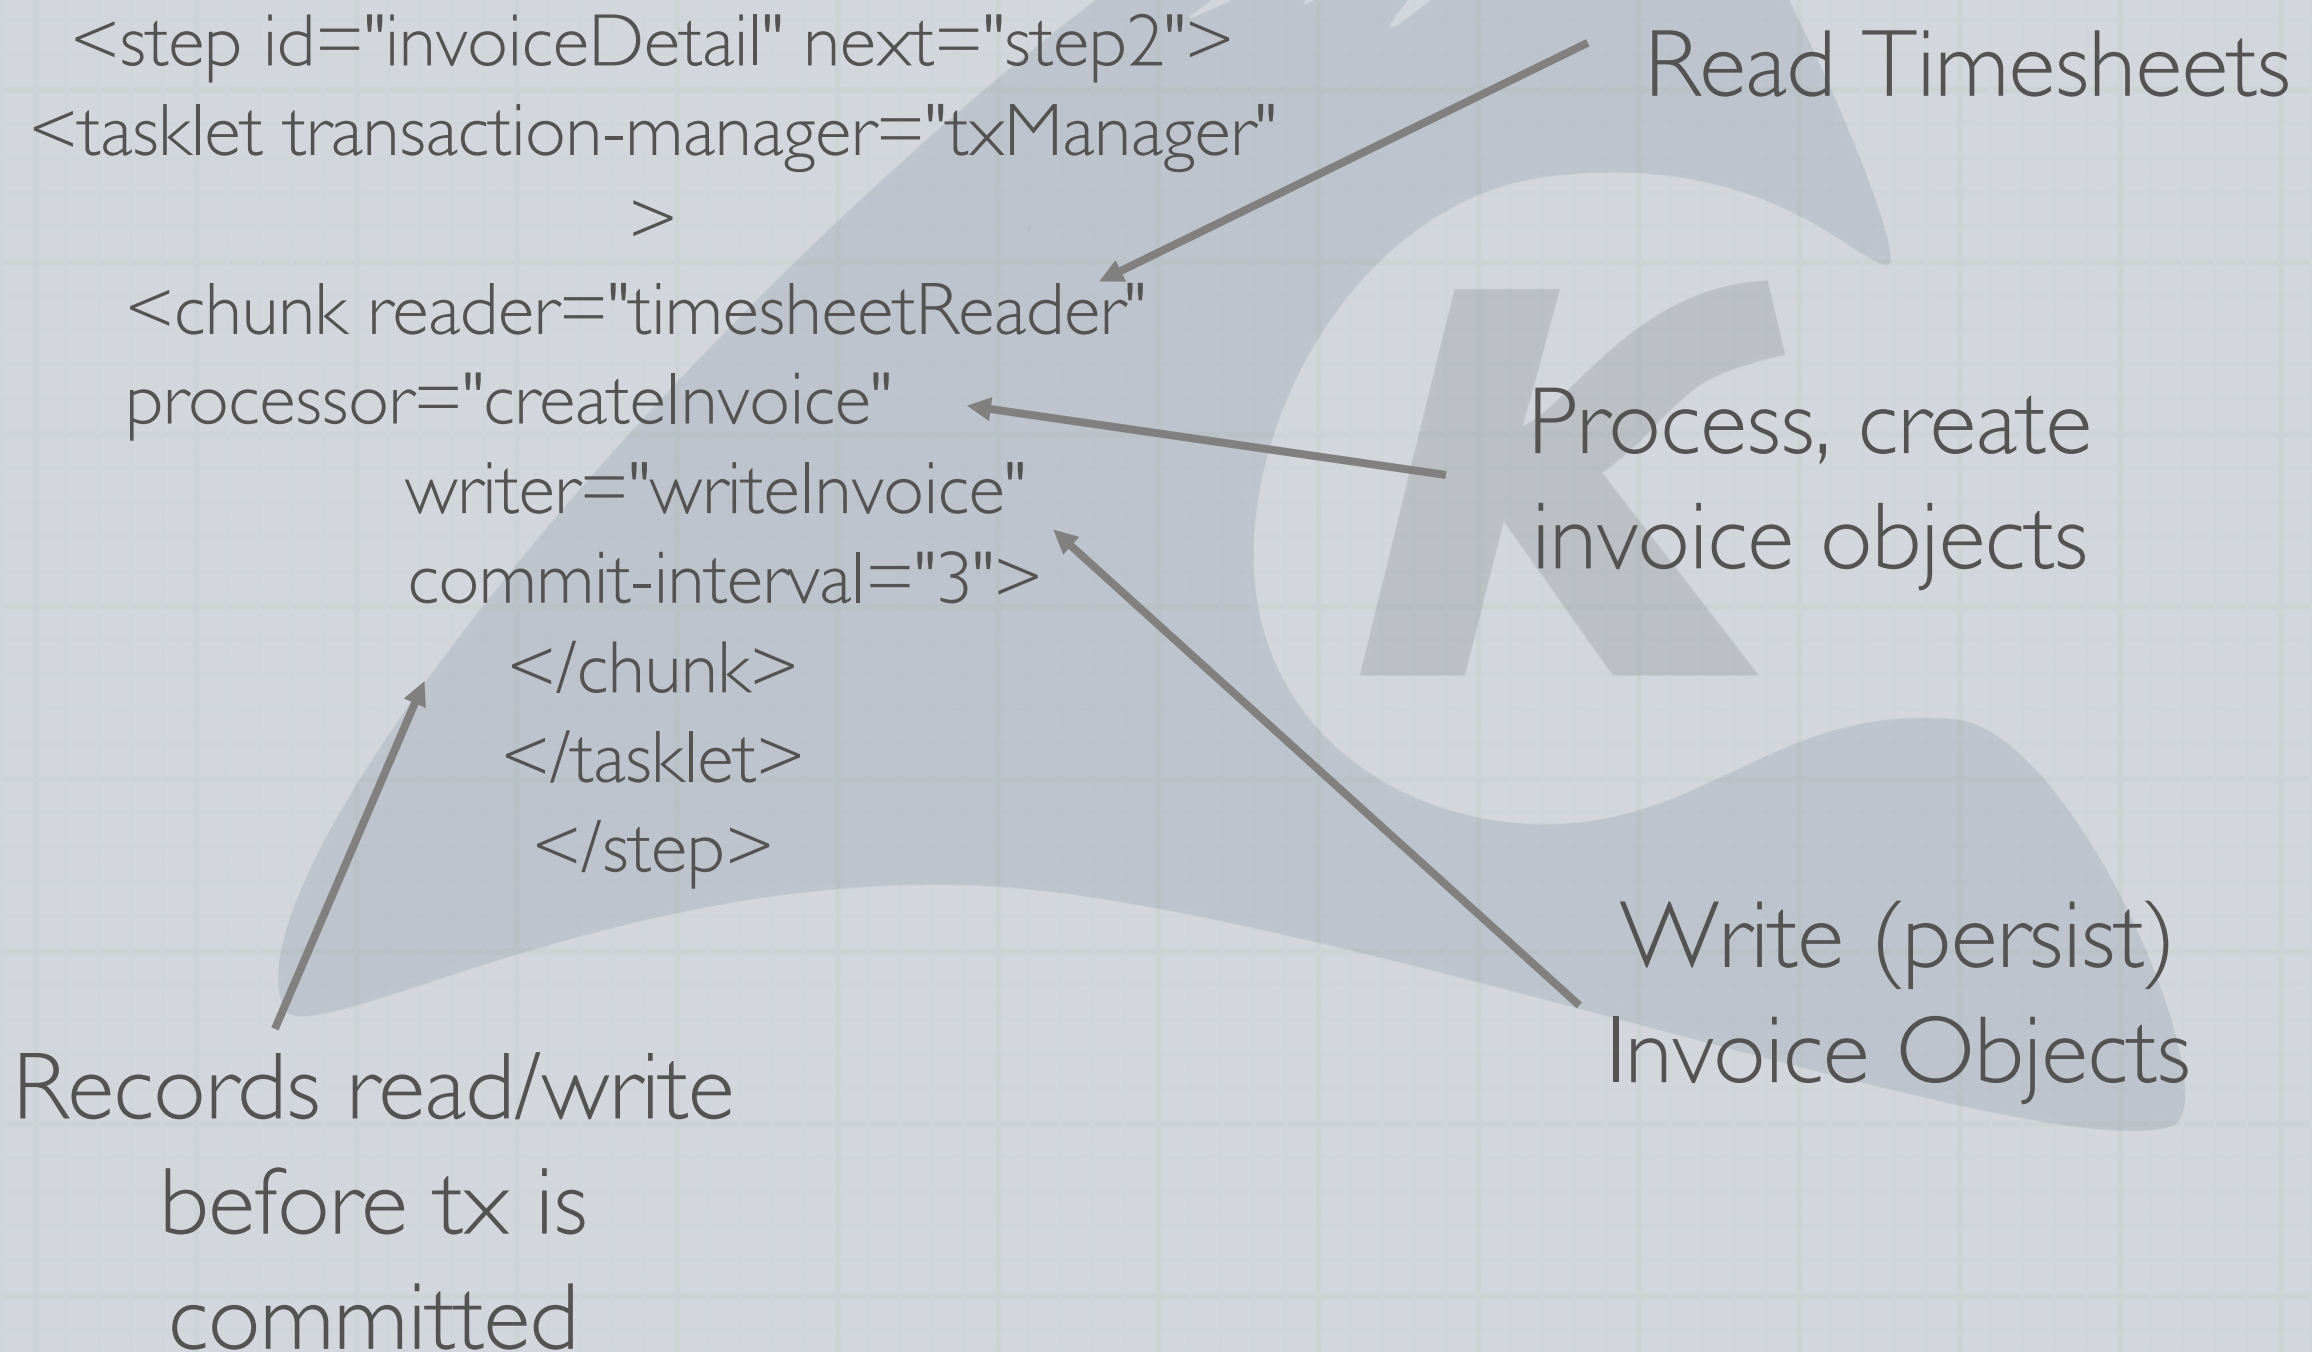
# Chunks - Where The Work Happens

- Tasklets
  - Contain Chunks
  - Chunks define Readers, Processors, Writer and Commit interval
  - Unit of Work
  - Allows HUGE amounts of data to be processed
- Developer implements Reader, Writer, Processor (optional)
  - Implements Interface in Spring Core classes



Chunk Elements

# Chunk Definition



# Supplied Readers/Writer

- Flat Files
- JMS
- Cursor and Paging (Hibernate, Spring, IBATIS, raw JDBC)
- XML
  - No coding required, just wire them up...
  - Many apply convention over configuration...
- HTTP/FTP

# Supplied JPA and FlatFile Reader

Read timesheet records from timesheet database...

```
<beans:bean id="timesheetReader"  
class="org.springframework.batch.item.database.JpaPagingItemReader">  
  <beans:property name="entityManagerFactory" ref="dataFactory"/>  
  <beans:property name="queryString" value="select t from Timesheet  
    t"/>  
  <beans:property name="pageSize" value="1000"/>  
</beans:bean>
```

# Supplied JPA and FlatFile Reader

Read timesheet records from comma delimited flat file...

```
<beans:bean id="timesheetFileItemReader"
  class="org.springframework.batch.item.file.FlatFileItemReader">
  <beans:property name="resource" value="classpath:data/input/timesheets.sdf" />
    <beans:property name="lineMapper">
      <beans:bean class="org.springframework.batch.item.file.mapping.DefaultLineMapper">
        <beans:property name="lineTokenizer">
          <beans:bean class="org.springframework.batch.item.file.transform.DelimitedLineTokenizer">
            <beans:property name="names"
              value="ID,week,company,monday,tuesday,wednesday,thursdsay,friday,saturday,sunday" />
          </beans:bean>
        </beans:property>
        <beans:property name="fieldSetMapper">
          <beans:bean class="com.khs.batch.timesheet.TimesheetFieldSetMapper" />
        </beans:property>
      </beans:bean>
    </beans:property>
  </beans:bean>
```



# FlatFileItemReader

## Java Config

```
@Bean
    public ItemReader<TickerData> reader() throws
        MalformedURLException {
        FlatFileItemReader<TickerData> reader = new
        FlatFileItemReader<TickerData>();
        reader.setResource(new
        UrlResource("http://finance.yahoo.com/d/quotes.csv?
        s=XOM+IBM+JNJ+MSFT&f=snd1o1p2"));
        reader.setLineMapper(new DefaultLineMapper<TickerData>() {{
            setLineTokenizer(new DelimitedLineTokenizer());
            setFieldSetMapper(new TickerFieldSetMapper());
        }});
        return reader;
    }
```



# Example - Reader Implementation

```
public class TimesheetReader implements ItemReader<Timesheet> {
```

```
List<Timesheet> timesheets;  
int current = 0;
```

```
@Autowired  
TimesheetJPADao dao;
```

```
@BeforeStep  
public void init() {  
    timesheets = dao.findAll();  
}
```

```
@AfterStep  
public void reset() {  
    timesheets = null;  
    current = 0;  
}
```

```
@Override  
public Timesheet read() throws Exception, UnexpectedInputException,  
ParseException, NonTransientResourceException {
```

```
    Timesheet ts = null;  
    if (current < timesheets.size()) {  
        ts = timesheets.get(current);  
        current++;  
    }  
    return ts;  
}
```

Initializing and  
Housekeeping  
Before/After Step  
Execution...

Read a  
Timesheet;  
Null Indicates  
Done...

# Example – Processor Implementation

Timesheet object  
from reader....

```
public class InvoiceProcessor implements ItemProcessor<Timesheet,  
InvoiceItem> {
```

```
@Override
```

```
public InvoiceItem process(Timesheet ts) throws Exception {
```

```
    InvoiceItem item = new InvoiceItem();
```

```
    item.setCompany(ts.getCompany());
```

```
    item.setAmount(ts.total() * 100.00);
```

```
    return item;
```

```
}  
}
```

InvoiceItem  
Object Goes  
to Writer...

# Example – Writer Implementation

```
public class TimesheetWriter implements ItemWriter<Timesheet> {
```

```
@Autowired
```

```
TimesheetJPADao dao;
```

```
@Override
```

```
public void write(List<? extends Timesheet> items) throws  
Exception {
```

```
    for (Timesheet ts : items) {  
        dao.persist(ts);
```

```
    }
```

```
}
```

```
}
```

# Intercepting Processing

- Listeners/Event Handler
  - Before Step/Job
  - After Step/Job
  - Job Completion



# Scalability Strategies

- Vertical
  - Faster machine, more RAM, CPU, GHz
- Horizontal
  - Add machines

# Spring Batch Scalability Support

- Remote Chunking
  - Chunks can be partitioned across process spaces
  - Using Messaging data read from chunk sent to a remote processor
- Step Partitioning
  - Step processing is partitioned to separate process or thread
  - Interacts with a partition event listener. JMS, Remote EJB, Web Service, Grid and more handlers are provided.



# Threads or Processes

- Multi-threaded Step (single process) – Vertical scaling  
Reader/Processor/Writer threaded
- Parallel Steps (single process) – Vertical scaling steps  
execute in parallel
- Remote Chunking of Step (multi process) – Horizontal  
scaling Processing/Writing on remote servers
- Partitioning a Step (single or multi process) –  
Horizontal scaling, Partitions reader data to multiple  
threads or remote nodes

# Example Parallel Step Execution

```
<job id="job1">
  <split id="split1" task-executor="taskExecutor" next="step4">
    <flow>
      <step id="step1" parent="s1" next="step2"/>
      <step id="step2" parent="s2"/>
    </flow>
    <flow>
      <step id="step3" parent="s3"/>
    </flow>
  </split>
  <step id="step4" parent="s4"/>
</job>

<beans:bean id="taskExecutor"
class="org.spr...SimpleAsyncTaskExecutor"/>
```

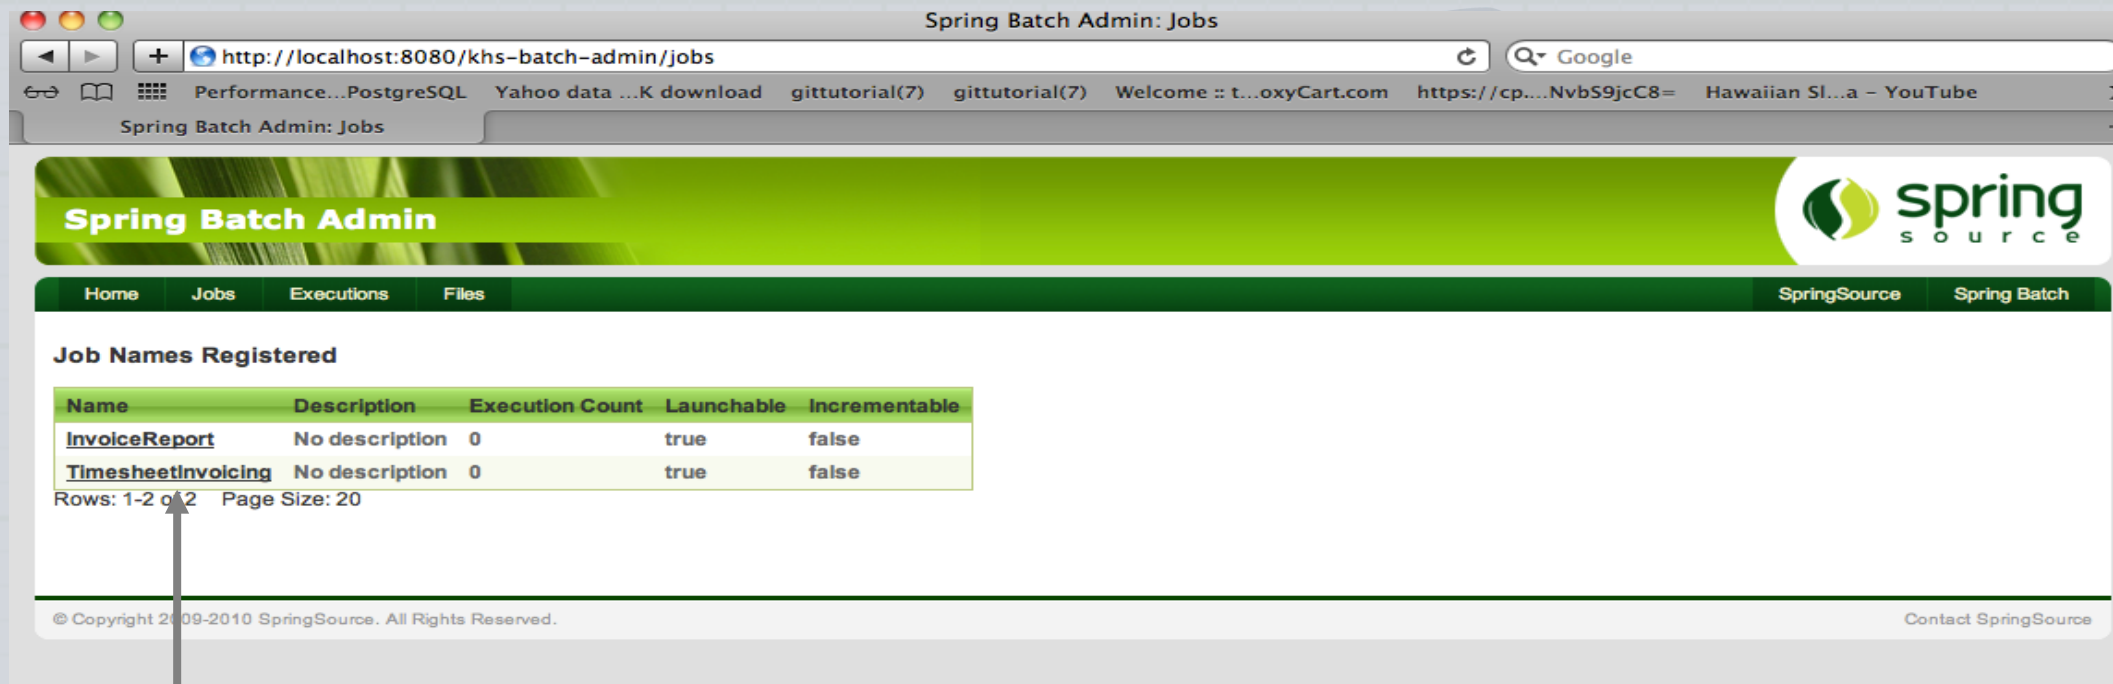
# Execution and Monitoring Environment

- Spring Batch Admin (Web App)
- Application Launchers (Command Line, JMS)
- OSGI Spring DM Server
- JMX (Used to interface with 3rd Party Scheduling and monitoring Software)
- Spring Integration (i.e. message bean or service to invoke job launcher(s))

# Scheduling

- Primitive support, does not really have full function scheduling features
- Out-of-the-Box
  - Quartz
  - Chron
- Third Party (JMX or custom API)

# Batch Admin Console



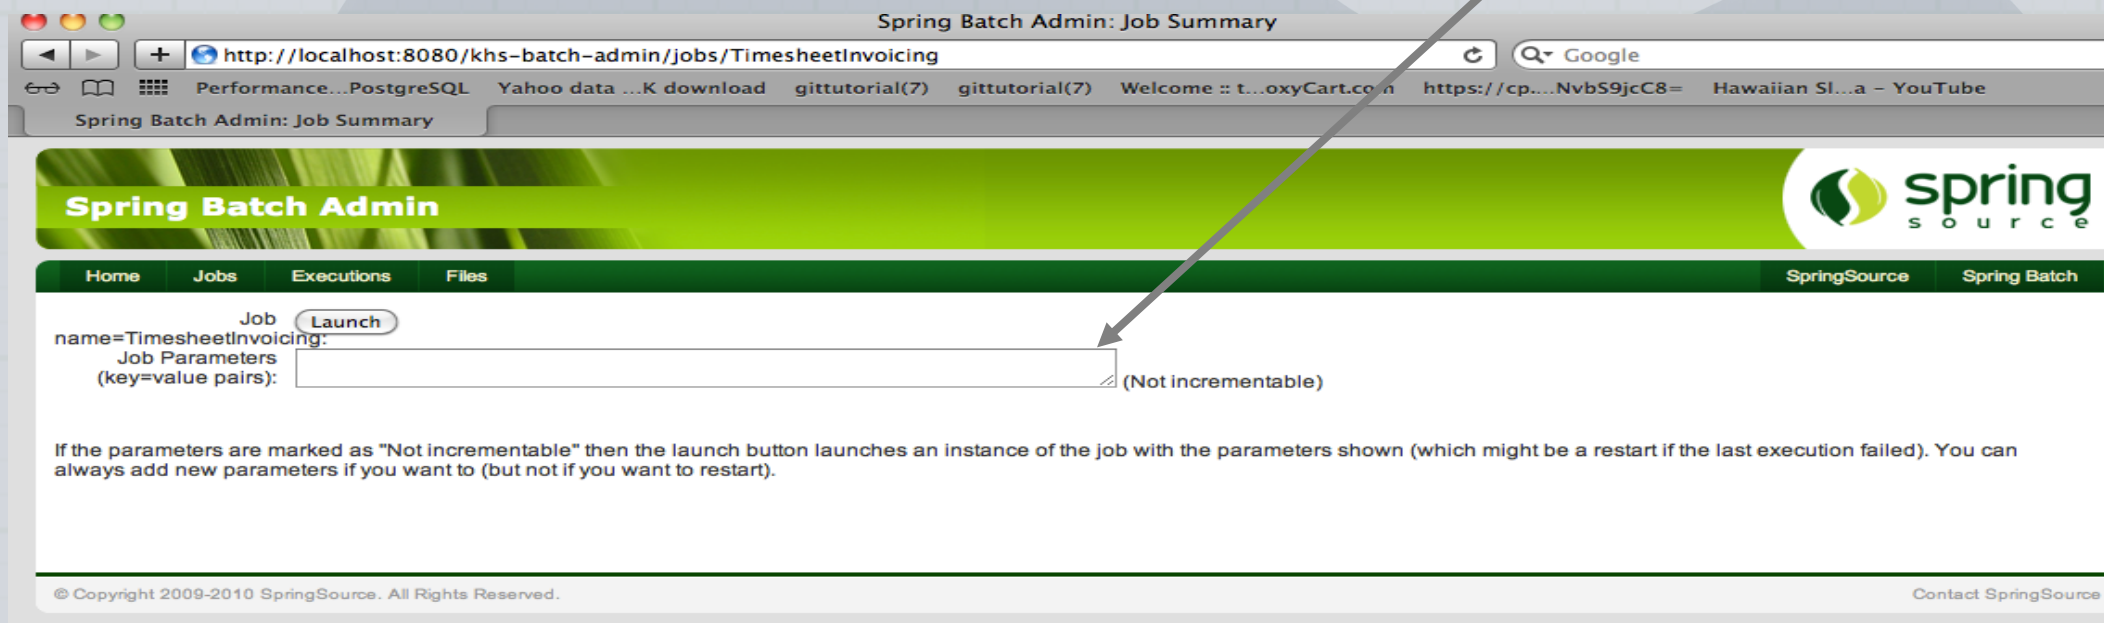
The screenshot shows the 'Spring Batch Admin: Jobs' page. The browser address bar is 'http://localhost:8080/khs-batch-admin/jobs'. The page has a green header with the 'Spring Batch Admin' title and the SpringSource logo. A navigation bar includes 'Home', 'Jobs', 'Executions', and 'Files'. The main content area is titled 'Job Names Registered' and contains a table with the following data:

| Name                               | Description    | Execution Count | Launchable | Incrementable |
|------------------------------------|----------------|-----------------|------------|---------------|
| <a href="#">InvoiceReport</a>      | No description | 0               | true       | false         |
| <a href="#">TimesheetInvoicing</a> | No description | 0               | true       | false         |

Below the table, it says 'Rows: 1-2 of 2' and 'Page Size: 20'. The footer includes '© Copyright 2009-2010 SpringSource. All Rights Reserved.' and a 'Contact SpringSource' link.

Execution Statistics

Job Launch  
Starting/Stopping



The screenshot shows the 'Spring Batch Admin: Job Summary' page for the 'TimesheetInvoicing' job. The browser address bar is 'http://localhost:8080/khs-batch-admin/jobs/TimesheetInvoicing'. The page has a green header with the 'Spring Batch Admin' title and the SpringSource logo. A navigation bar includes 'Home', 'Jobs', 'Executions', and 'Files'. The main content area shows the job name 'TimesheetInvoicing' and a 'Launch' button. Below the button is a text input field for 'Job Parameters (key=value pairs):'. To the right of the input field, it says '(Not incrementable)'. Below the input field, there is a paragraph of text: 'If the parameters are marked as "Not incrementable" then the launch button launches an instance of the job with the parameters shown (which might be a restart if the last execution failed). You can always add new parameters if you want to (but not if you want to restart).' The footer includes '© Copyright 2009-2010 SpringSource. All Rights Reserved.' and a 'Contact SpringSource' link.



# Now, A Demo!

For more, check out our 6-part tutorial series on how to get started with Spring Batch on <https://keyholesoftware.com/blog>.

**Short Link: [bit.ly/springbatchkhs](https://bit.ly/springbatchkhs)**